

Forest Image Segmentation - UNet

```
In [9]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import os
import random
import cv2
import tensorflow as tf
from tensorflow.python.keras.models import Sequential
from tensorflow.keras.preprocessing.image import load_img, img_to_array
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
from tensorflow.keras.layers import Conv2D, BatchNormalization, Activation, MaxPool
from tensorflow.keras.models import Model
```

Directory Setup

```
In [11]: base_directory = r'D:\data science\Practice\Unet\Forest Segmented\Forest Segmented'
images_folder = os.path.join(base_directory, 'image')
masks_folder = os.path.join(base_directory, 'mask')
data = pd.read_csv(os.path.join(base_directory, 'meta_data.csv'))
```

Data Augmentation

```
In [30]: img_dim = 256

image_datagen = ImageDataGenerator(rescale=1./255, validation_split=0.15)
mask_datagen = ImageDataGenerator(rescale=1./255, validation_split=0.15)

train_image_generator = image_datagen.flow_from_directory(
    r'D:\data science\Practice\Unet\Forest Segmented\Forest Segmented',
    target_size=(img_dim, img_dim),
    class_mode = None,
    classes = ['image'],
    batch_size = 8,
    seed=42,
    subset='training')

train_mask_generator = mask_datagen.flow_from_directory(
    r'D:\data science\Practice\Unet\Forest Segmented\Forest Segmented',
    target_size=(img_dim, img_dim),
    class_mode = None,
    classes = ['mask'],
    color_mode = 'grayscale',
    batch_size = 8,
    seed=42,
    subset='training')
```

```

val_image_generator = image_datagen.flow_from_directory(
    r'D:\data science\Practice\Unet\Forest Segmented\Forest Segmented',
    target_size=(img_dim, img_dim),
    class_mode = None,
    classes = ['image'],
    batch_size = 8,
    seed=42,
    subset='validation')

val_mask_generator = mask_datagen.flow_from_directory(
    r'D:\data science\Practice\Unet\Forest Segmented\Forest Segmented',
    target_size=(img_dim, img_dim),
    class_mode = None,
    classes = ['mask'],
    color_mode = 'grayscale',
    batch_size = 8,
    seed=42,
    subset='validation')

train_generator = zip(train_image_generator, train_mask_generator)
val_generator = zip(val_image_generator, val_mask_generator)

```

Found 4342 images belonging to 1 classes.
 Found 4342 images belonging to 1 classes.
 Found 766 images belonging to 1 classes.
 Found 766 images belonging to 1 classes.

Display Sample Training Images

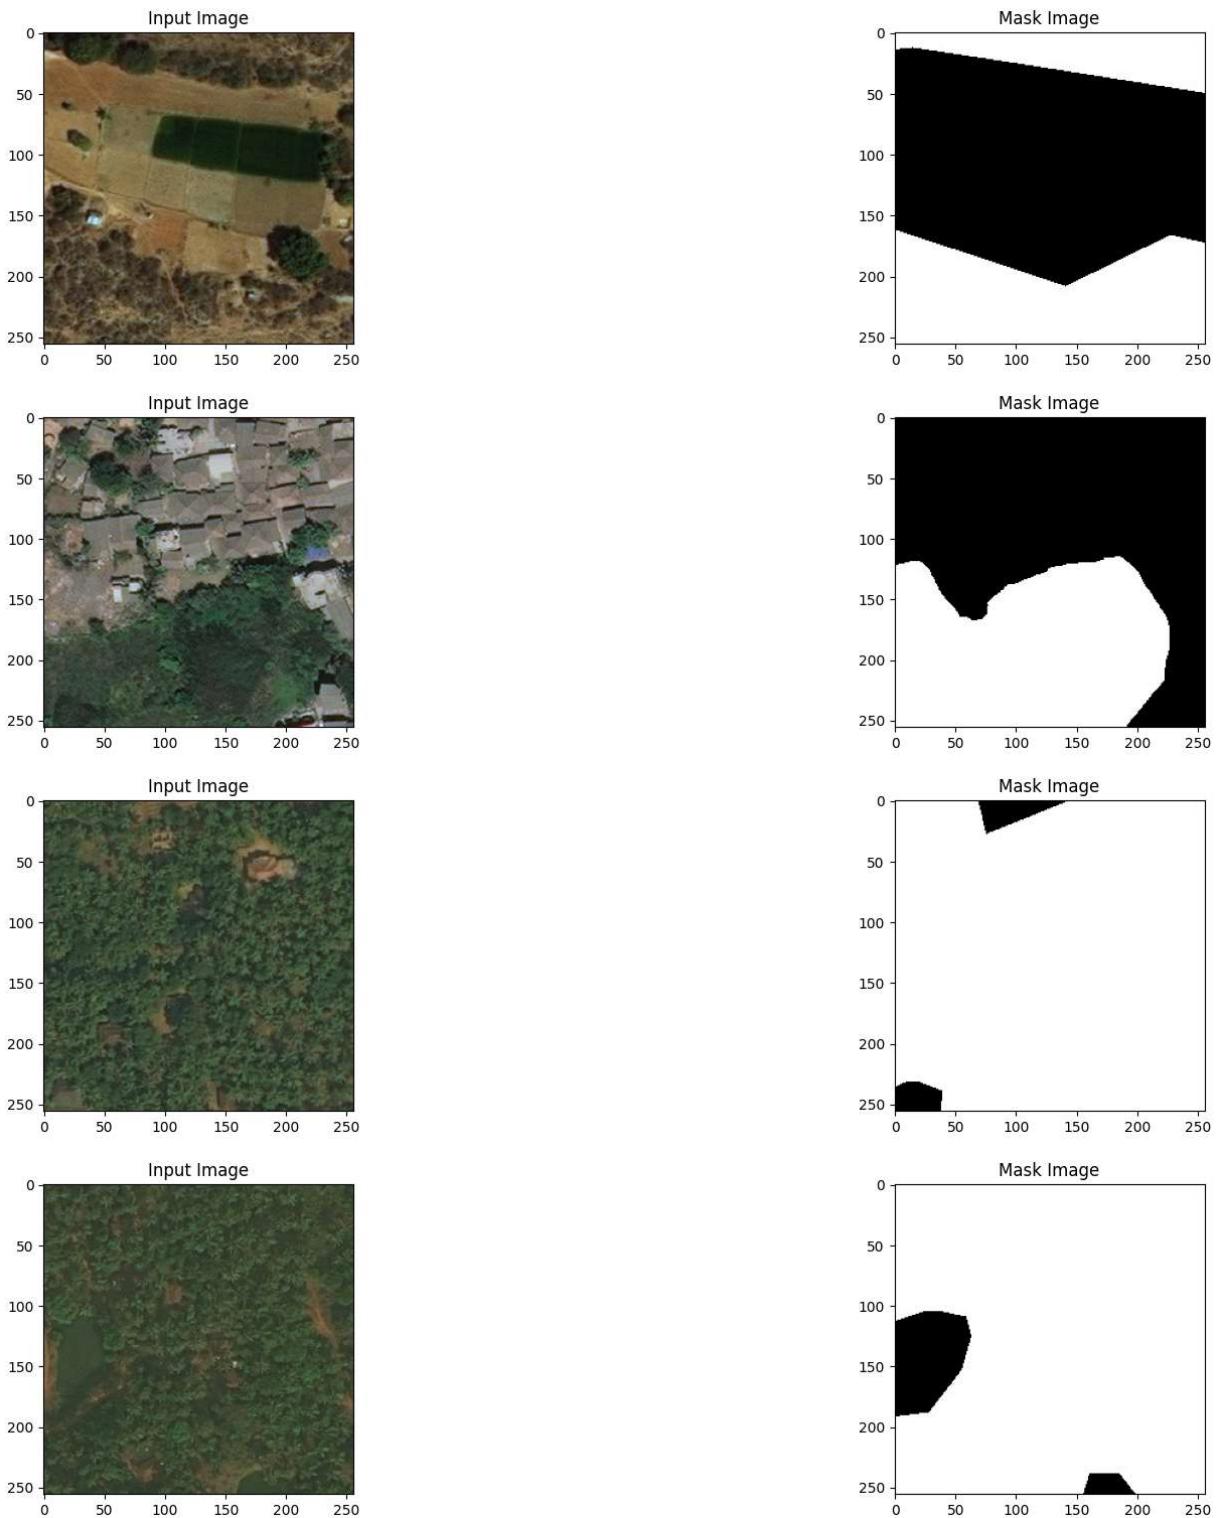
In [31]:

```
training_samples_size = train_image_generator.samples
val_samples_size = val_image_generator.samples
```

In [32]:

```
n = 0
for i,m in train_generator:
    img,mask = i,m

    if n < 5:
        fig, axs = plt.subplots(1 , 2, figsize=(20,4))
        axs[0].imshow(img[0])
        axs[0].set_title('Input Image')
        axs[1].imshow(mask[0],cmap='gray')
        axs[1].set_title('Mask Image')
        plt.show()
        n+=1
    else:
        break
```





Unet Model

```
In [33]: def build_unet(input_shape):
    def conv_block(input, num_filters):
        x = Conv2D(num_filters, 3, padding="same")(input)
        x = BatchNormalization()(x)
        x = Activation("relu")(x)

        x = Conv2D(num_filters, 3, padding="same")(x)
        x = BatchNormalization()(x)
        x = Activation("relu")(x)

    return x

    def encoder_block(input, num_filters):
        x = conv_block(input, num_filters)
        p = MaxPool2D((2, 2))(x)
        return x, p

    def decoder_block(input, skip_features, num_filters):
        x = Conv2DTranspose(num_filters, (2, 2), strides=2, padding="same")(input)
        x = Concatenate()([x, skip_features])
        x = conv_block(x, num_filters)
        return x

    inputs = Input(input_shape)

    s1, p1 = encoder_block(inputs, 32)
    s2, p2 = encoder_block(p1, 64)
    s3, p3 = encoder_block(p2, 128)
    s4, p4 = encoder_block(p3, 256)

    b1 = conv_block(p4, 512)

    d1 = decoder_block(b1, s4, 256)
    d2 = decoder_block(d1, s3, 128)
    d3 = decoder_block(d2, s2, 64)
    d4 = decoder_block(d3, s1, 32)

    outputs = Conv2D(1, 1, padding="same", activation="sigmoid")(d4)
```

```
model = Model(inputs, outputs, name="U-Net")
return model
```

```
In [34]: input_shape = (img_dim, img_dim, 3)
model = build_unet(input_shape)
model.compile(optimizer = tf.keras.optimizers.Adam(lr = 0.001), loss = ['binary_cro
model.summary()
```

Model: "U-Net"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_4 (InputLayer)	[(None, 256, 256, 3 0)]		[]
conv2d_57 (Conv2D)	(None, 256, 256, 32 896)		['input_4[0][0]']
batch_normalization_54 (BatchN ormalization)	(None, 256, 256, 32 128)		['conv2d_57[0][0]']
activation_54 (Activation)	(None, 256, 256, 32 0 on_54[0][0]')		['batch_normalizati
conv2d_58 (Conv2D)	(None, 256, 256, 32 9248 [0]')		['activation_54[0]
batch_normalization_55 (BatchN ormalization)	(None, 256, 256, 32 128)		['conv2d_58[0][0]']
activation_55 (Activation)	(None, 256, 256, 32 0 on_55[0][0]')		['batch_normalizati
max_pooling2d_12 (MaxPooling2D)	(None, 128, 128, 32 0 [0]')		['activation_55[0]
conv2d_59 (Conv2D)	(None, 128, 128, 64 18496 [0][0]')		['max_pooling2d_12
batch_normalization_56 (BatchN ormalization)	(None, 128, 128, 64 256)		['conv2d_59[0][0]']
activation_56 (Activation)	(None, 128, 128, 64 0 on_56[0][0]')		['batch_normalizati
conv2d_60 (Conv2D)	(None, 128, 128, 64 36928 [0]')		['activation_56[0]
batch_normalization_57 (BatchN ormalization)	(None, 128, 128, 64 256)		['conv2d_60[0][0]']
activation_57 (Activation)	(None, 128, 128, 64 0 on_57[0][0]')		['batch_normalizati

max_pooling2d_13 (MaxPooling2D [0])	(None, 64, 64, 64) 0	['activation_57[0]']
conv2d_61 (Conv2D [0][0])	(None, 64, 64, 128) 73856	['max_pooling2d_13 [0][0]']
batch_normalization_58 (BatchN ormalization)	(None, 64, 64, 128) 512	['conv2d_61[0][0]']
activation_58 (Activation on_58[0][0])	(None, 64, 64, 128) 0	['batch_norma lization_58[0][0]']
conv2d_62 (Conv2D [0])	(None, 64, 64, 128) 147584	['activation_58[0] [0]']
batch_normalization_59 (BatchN ormalization)	(None, 64, 64, 128) 512	['conv2d_62[0][0]']
activation_59 (Activation on_59[0][0])	(None, 64, 64, 128) 0	['batch_norma lization_59[0][0]']
max_pooling2d_14 (MaxPooling2D [0])	(None, 32, 32, 128) 0	['activation_59[0] [0]']
conv2d_63 (Conv2D [0][0])	(None, 32, 32, 256) 295168	['max_pooling2d_14 [0][0]']
batch_normalization_60 (BatchN ormalization)	(None, 32, 32, 256) 1024	['conv2d_63[0][0]']
activation_60 (Activation on_60[0][0])	(None, 32, 32, 256) 0	['batch_norma lization_60[0][0]']
conv2d_64 (Conv2D [0])	(None, 32, 32, 256) 590080	['activation_60[0] [0]']
batch_normalization_61 (BatchN ormalization)	(None, 32, 32, 256) 1024	['conv2d_64[0][0]']
activation_61 (Activation on_61[0][0])	(None, 32, 32, 256) 0	['batch_norma lization_61[0][0]']
max_pooling2d_15 (MaxPooling2D [0])	(None, 16, 16, 256) 0	['activation_61[0] [0]']
conv2d_65 (Conv2D [0][0])	(None, 16, 16, 512) 1180160	['max_pooling2d_15 [0][0]']
batch_normalization_62 (BatchN ormalization)	(None, 16, 16, 512) 2048	['conv2d_65[0][0]']
activation_62 (Activation on_62[0][0])	(None, 16, 16, 512) 0	['batch_norma lization_62[0][0]']

conv2d_66 (Conv2D) [0]'	(None, 16, 16, 512)	2359808	['activation_62[0]
batch_normalization_63 (BatchN ormalization)	(None, 16, 16, 512)	2048	['conv2d_66[0][0]']
activation_63 (Activation) on_63[0][0]']	(None, 16, 16, 512)	0	['batch_norma
conv2d_transpose_12 (Conv2DTra nspose)	(None, 32, 32, 256)	524544	['activation_63[0]
concatenate_12 (Concatenate) 12[0][0]', [0]']	(None, 32, 32, 512)	0	['conv2d_transpose_
activation_61 (Activation) on_61[0]']			'activation_61[0]
conv2d_67 (Conv2D) [0]'	(None, 32, 32, 256)	1179904	['concatenate_12[0]
batch_normalization_64 (BatchN ormalization)	(None, 32, 32, 256)	1024	['conv2d_67[0][0]']
activation_64 (Activation) on_64[0][0]']	(None, 32, 32, 256)	0	['batch_norma
conv2d_68 (Conv2D) [0]']	(None, 32, 32, 256)	590080	['activation_64[0]
batch_normalization_65 (BatchN ormalization)	(None, 32, 32, 256)	1024	['conv2d_68[0][0]']
activation_65 (Activation) on_65[0][0]']	(None, 32, 32, 256)	0	['batch_norma
conv2d_transpose_13 (Conv2DTra nspose)	(None, 64, 64, 128)	131200	['activation_65[0]
concatenate_13 (Concatenate) 13[0][0]', [0]']	(None, 64, 64, 256)	0	['conv2d_transpose_
activation_59 (Activation) on_59[0]']			'activation_59[0]
conv2d_69 (Conv2D) [0]'	(None, 64, 64, 128)	295040	['concatenate_13[0]
batch_normalization_66 (BatchN ormalization)	(None, 64, 64, 128)	512	['conv2d_69[0][0]']
activation_66 (Activation) on_66[0][0]']	(None, 64, 64, 128)	0	['batch_norma
conv2d_70 (Conv2D)	(None, 64, 64, 128)	147584	['activation_66[0]

```
[0]']

batch_normalization_67 (BatchN (None, 64, 64, 128) 512      ['conv2d_70[0][0]']
ormalization)

activation_67 (Activation)      (None, 64, 64, 128) 0      ['batch_normalizati
on_67[0][0]']

conv2d_transpose_14 (Conv2DTra (None, 128, 128, 64 32832  ['activation_67[0]
[0]']
nspose)                      )

concatenate_14 (Concatenate)   (None, 128, 128, 12 0      ['conv2d_transpose_
14[0][0]',

[0]']]

conv2d_71 (Conv2D)            (None, 128, 128, 64 73792  ['concatenate_14[0]
[0]']
)

batch_normalization_68 (BatchN (None, 128, 128, 64 256      ['conv2d_71[0][0]']
ormalization)
)

activation_68 (Activation)    (None, 128, 128, 64 0      ['batch_normalizati
on_68[0][0]']
)

conv2d_72 (Conv2D)            (None, 128, 128, 64 36928  ['activation_68[0]
[0]']
)

batch_normalization_69 (BatchN (None, 128, 128, 64 256      ['conv2d_72[0][0]']
ormalization)
)

activation_69 (Activation)    (None, 128, 128, 64 0      ['batch_normalizati
on_69[0][0]']
)

conv2d_transpose_15 (Conv2DTra (None, 256, 256, 32 8224  ['activation_69[0]
[0]']
nspose)
)

concatenate_15 (Concatenate)  (None, 256, 256, 64 0      ['conv2d_transpose_
15[0][0]',

[0]'])

conv2d_73 (Conv2D)            (None, 256, 256, 32 18464  ['concatenate_15[0]
[0]']
)

batch_normalization_70 (BatchN (None, 256, 256, 32 128      ['conv2d_73[0][0]']
ormalization)
)

activation_70 (Activation)    (None, 256, 256, 32 0      ['batch_normalizati
```

```
on_70[0][0]')
)
conv2d_74 (Conv2D)      (None, 256, 256, 32  9248      ['activation_70[0]
[0]')
)
batch_normalization_71 (BatchN  (None, 256, 256, 32  128      ['conv2d_74[0][0]')
ormalization)
)
activation_71 (Activation)   (None, 256, 256, 32  0      ['batch_normalizati
on_71[0][0]')
)
conv2d_75 (Conv2D)        (None, 256, 256, 1)  33      ['activation_71[0]
[0]')
=====
=====
Total params: 7,771,873
Trainable params: 7,765,985
Non-trainable params: 5,888
```

Model Training

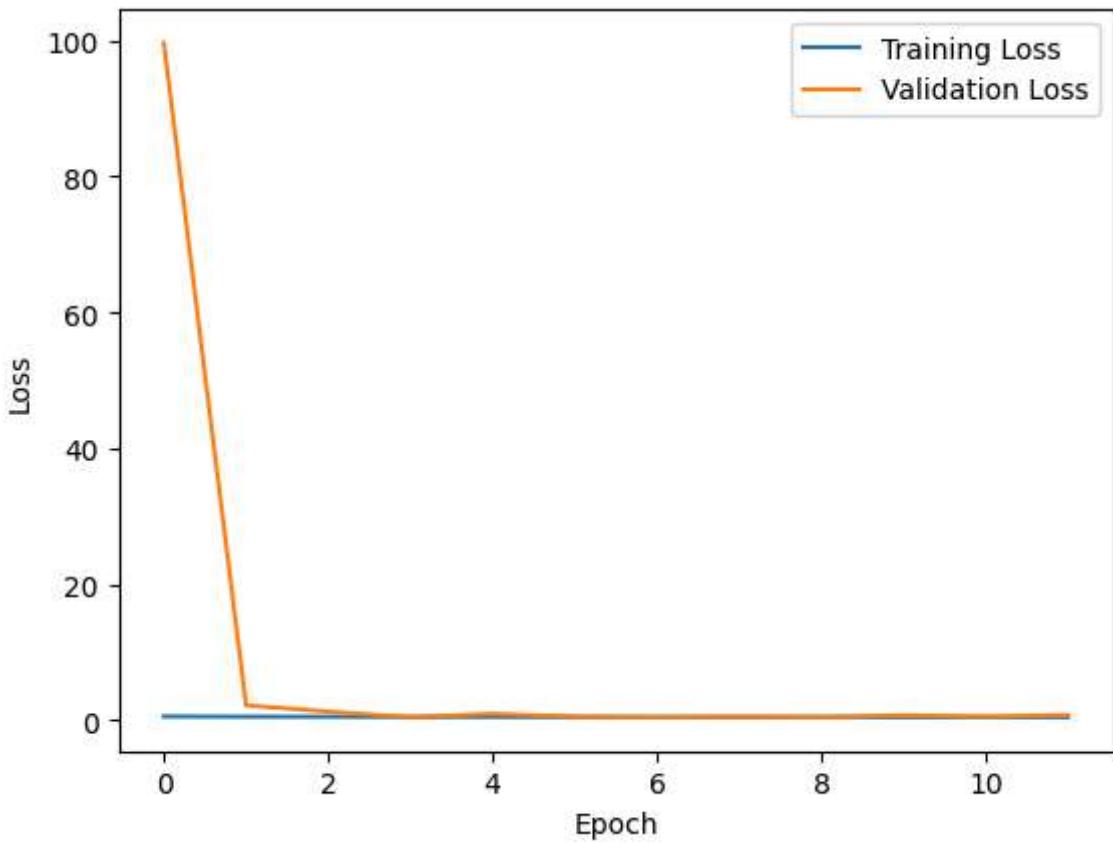
```
In [35]: earlystop = EarlyStopping(monitor = 'val_loss',
                                    min_delta = 0,
                                    patience = 5,
                                    verbose = 1,
                                    restore_best_weights = True)

history = model.fit(train_generator,
                     steps_per_epoch=training_samples_size//32,
                     validation_data=val_generator,
                     validation_steps=val_samples_size//32,
                     epochs=50, callbacks=[earlystop])
```

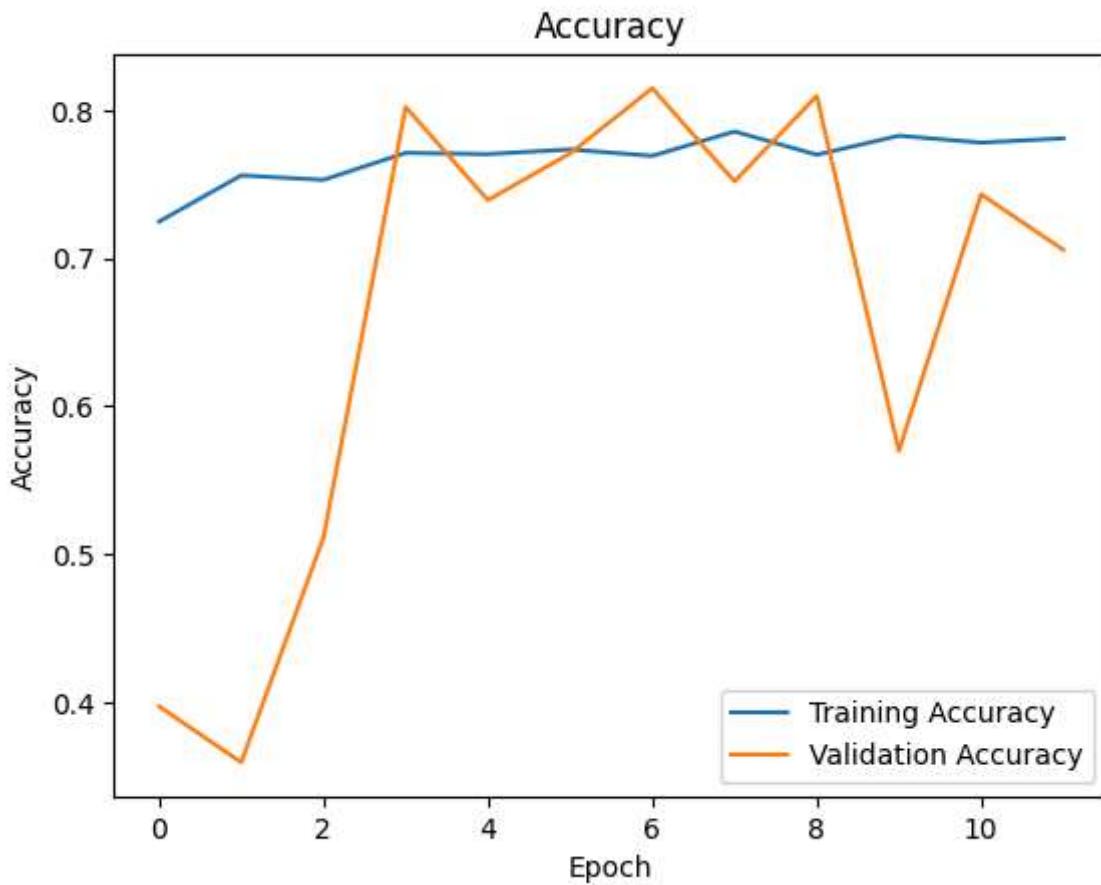
```
Epoch 1/50
135/135 [=====] - 54s 354ms/step - loss: 0.5592 - accuracy: 0.7247 - val_loss: 99.6630 - val_accuracy: 0.3975
Epoch 2/50
135/135 [=====] - 47s 351ms/step - loss: 0.5199 - accuracy: 0.7559 - val_loss: 2.1873 - val_accuracy: 0.3597
Epoch 3/50
135/135 [=====] - 48s 352ms/step - loss: 0.5173 - accuracy: 0.7529 - val_loss: 1.2753 - val_accuracy: 0.5116
Epoch 4/50
135/135 [=====] - 51s 376ms/step - loss: 0.4886 - accuracy: 0.7711 - val_loss: 0.4415 - val_accuracy: 0.8021
Epoch 5/50
135/135 [=====] - 48s 353ms/step - loss: 0.4952 - accuracy: 0.7703 - val_loss: 0.9443 - val_accuracy: 0.7393
Epoch 6/50
135/135 [=====] - 48s 353ms/step - loss: 0.4788 - accuracy: 0.7734 - val_loss: 0.5278 - val_accuracy: 0.7707
Epoch 7/50
135/135 [=====] - 48s 354ms/step - loss: 0.4873 - accuracy: 0.7691 - val_loss: 0.4221 - val_accuracy: 0.8148
Epoch 8/50
135/135 [=====] - 48s 353ms/step - loss: 0.4586 - accuracy: 0.7856 - val_loss: 0.5412 - val_accuracy: 0.7518
Epoch 9/50
135/135 [=====] - 47s 345ms/step - loss: 0.4879 - accuracy: 0.7699 - val_loss: 0.4353 - val_accuracy: 0.8095
Epoch 10/50
135/135 [=====] - 47s 349ms/step - loss: 0.4616 - accuracy: 0.7827 - val_loss: 0.7347 - val_accuracy: 0.5700
Epoch 11/50
135/135 [=====] - 48s 354ms/step - loss: 0.4648 - accuracy: 0.7781 - val_loss: 0.5161 - val_accuracy: 0.7432
Epoch 12/50
135/135 [=====] - ETA: 0s - loss: 0.4684 - accuracy: 0.7809
Restoring model weights from the end of the best epoch: 7.
135/135 [=====] - 48s 355ms/step - loss: 0.4684 - accuracy: 0.7809 - val_loss: 0.7903 - val_accuracy: 0.7057
Epoch 12: early stopping
```

Results

```
In [36]: plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [37]: plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Accuracy')
plt.show()
```



Testing

```
In [41]: from skimage.io import imread
        from skimage.transform import resize

# Function to predict segmentation for a custom image
def predict_segmentation(model, custom_image_path):
    custom_image = imread(custom_image_path)
    custom_image = resize(custom_image, (256, 256), preserve_range=True) # Resize
    custom_image = custom_image.astype('float32') / 255.0
    custom_image = np.expand_dims(custom_image, axis=0)
    predicted_mask = model.predict(custom_image)
    return predicted_mask[0]

# Path to the custom image
custom_image_path = r"C:\Users\deepchanddc2\Downloads\images.jpg"

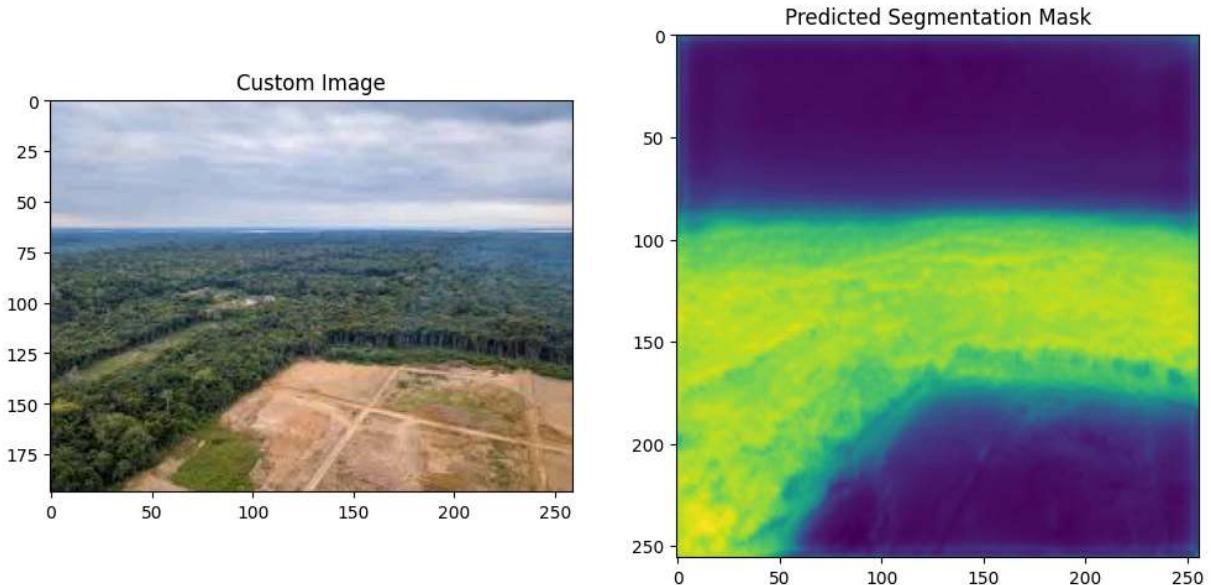
# Predict segmentation for the custom image
predicted_mask = predict_segmentation(model, custom_image_path)

# Visualize the custom image and its predicted mask
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.imshow(imread(custom_image_path))
plt.title('Custom Image')
```

```
plt.subplot(1, 2, 2)
plt.imshow(predicted_mask[:, :, 0], cmap='viridis')
plt.title('Predicted Segmentation Mask')

plt.show()
```

1/1 [=====] - 0s 21ms/step



```
In [43]: model.save(r'D:\data science\Practice\Unet\Model')

print("Model saved successfully")
```

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 23). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: D:\data science\Practice\Unet\Model\assets

INFO:tensorflow:Assets written to: D:\data science\Practice\Unet\Model\assets

Model saved successfully